Page 1: Introduction to RAG Systems

Retrieval-Augmented Generation combines external document retrieval with LLMs.

This allows AI models to answer questions using private knowledge bases accurately.

Page 2: Ingestion Pipeline Architecture

Documents are parsed into clean text and sliced into overlapping chunks.

Each chunk is transformed into vector embeddings and saved into the vector store.

Page 3: Query and Retrieval Workflow

User queries are converted into vectors and matched using cosine similarity.

Top ranked chunks are injected into the prompt as grounded reference context.